



BEUTH HOCHSCHULE FÜR TECHNIK BERLIN
University of Applied Sciences

Learning from Images

Image Descriptors and Marker detection

Master DataScience
Winter term 2019/20

Prof. Dr. Kristian Hildebrand
khildebrand@beuth-hochschule.de

Course outline

- **Image descriptors**
- Image matching



What do you see?

**One step back:
How do we describe an image?
How do we compare images?**

Image descriptors

- Descriptor - a piece of stored information that is used to identify an item in an information storage and retrieval system.
 - Descriptors save pertinent information, saving the processing time in future queries
 - Descriptors can save image features that are essential for search and comparison

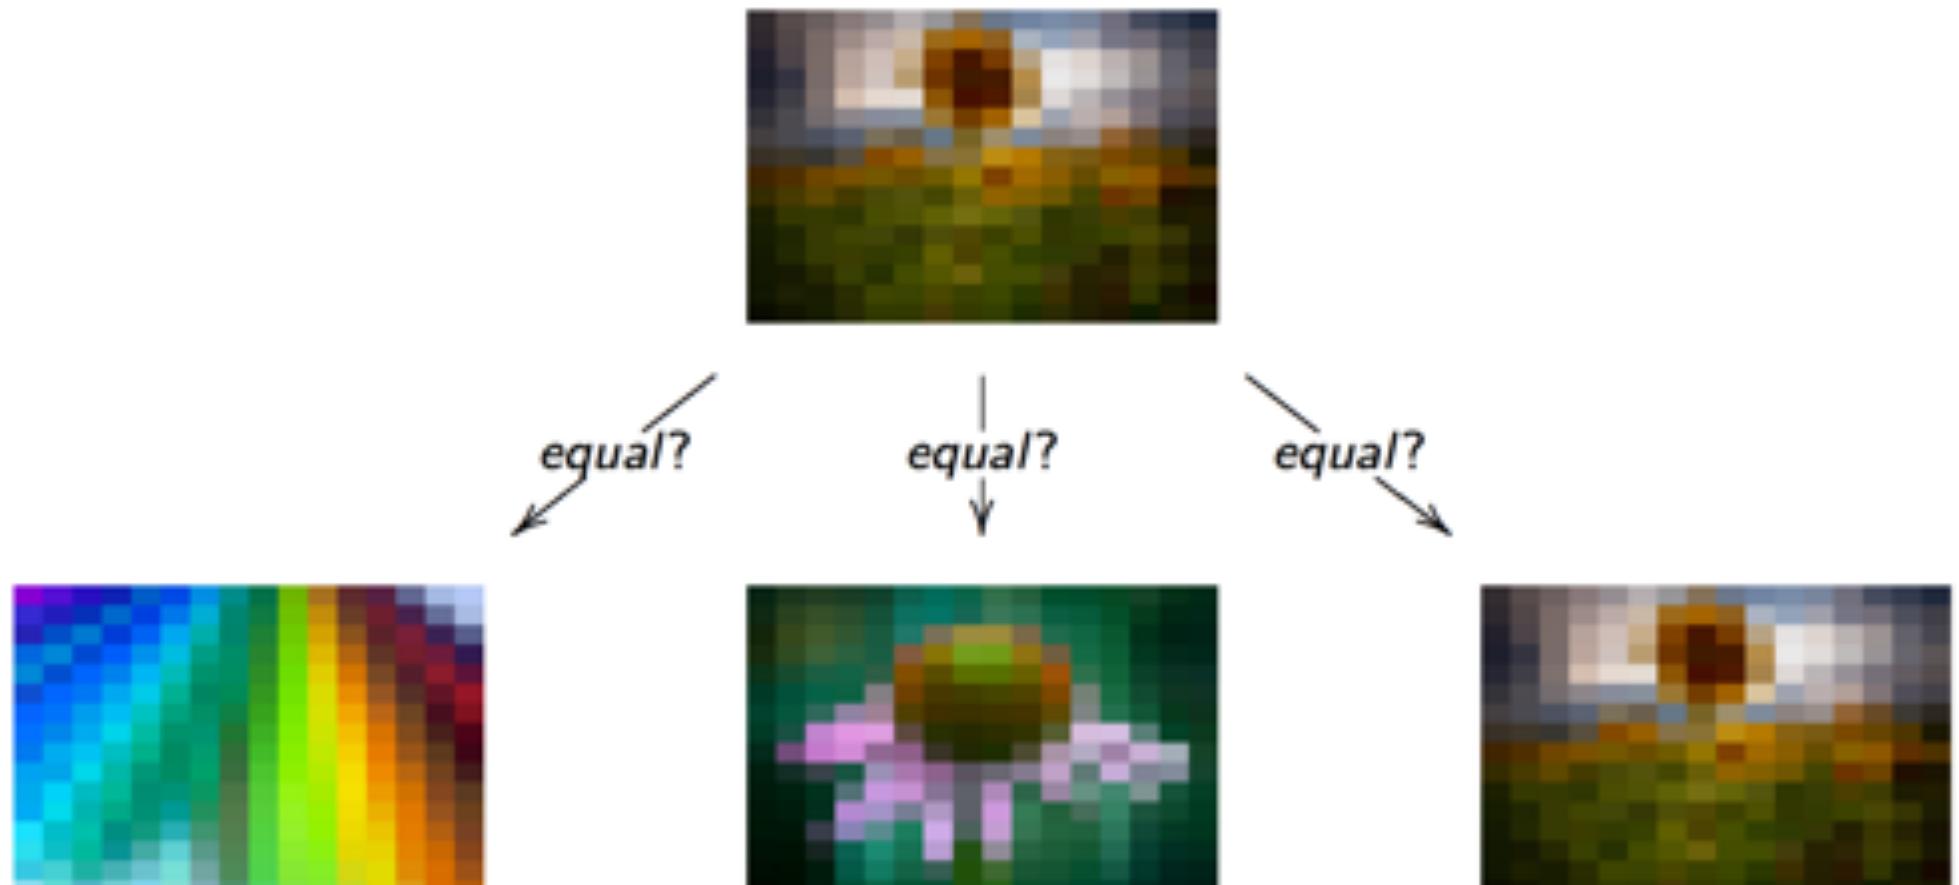
Global image descriptors



<https://groups.csail.mit.edu/vision/TinyImages/>

<http://www.flickr.com/photos/alphageek/2759567956/>

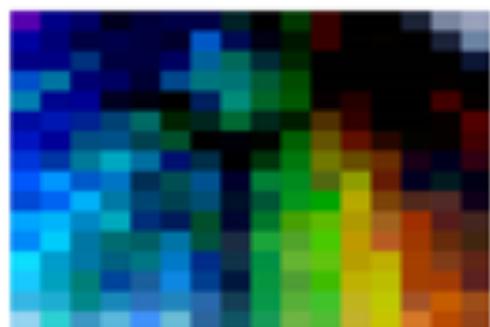
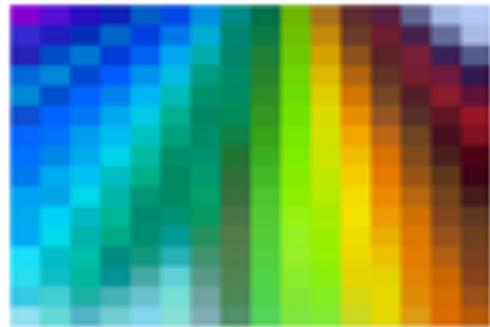
Global image descriptors



<http://www.flickr.com/photos/40513596@N00/71762740>

<http://www.flickr.com/photos/alphageek/2993813155/>

Comparison example



14862080

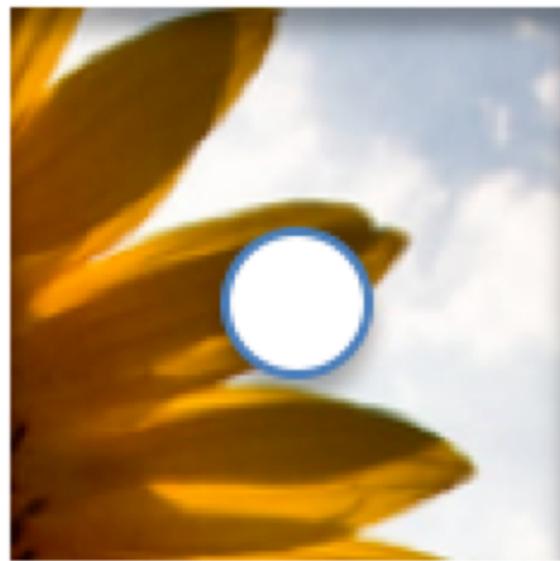
437120

0

Global vs Local Features

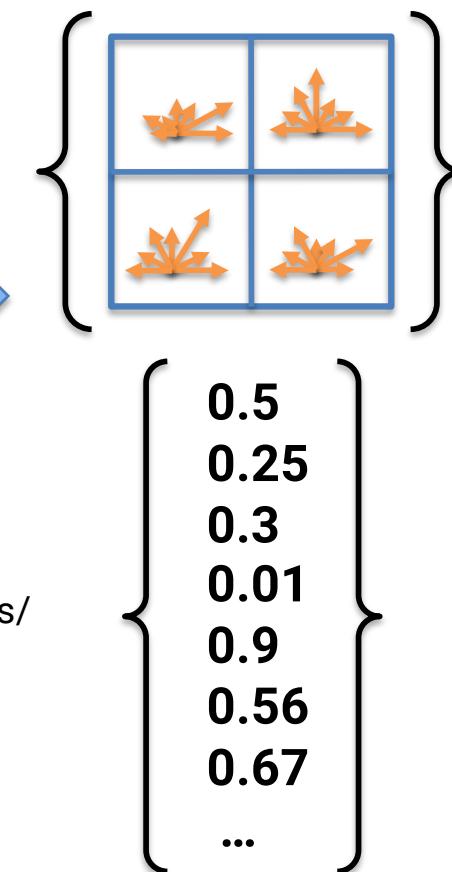


Image as a vector



Feature
description

Histogram:
<http://tinlizzie.org/histograms/>



Why local features?

- Locality
 - features are local, so robust to occlusion and clutter
- Distinctiveness:
 - can differentiate a large database of objects
- Quantity
 - hundreds or thousands in a single image
- Efficiency
 - real-time performance achievable
- Generality
 - exploit different types of features in different situations

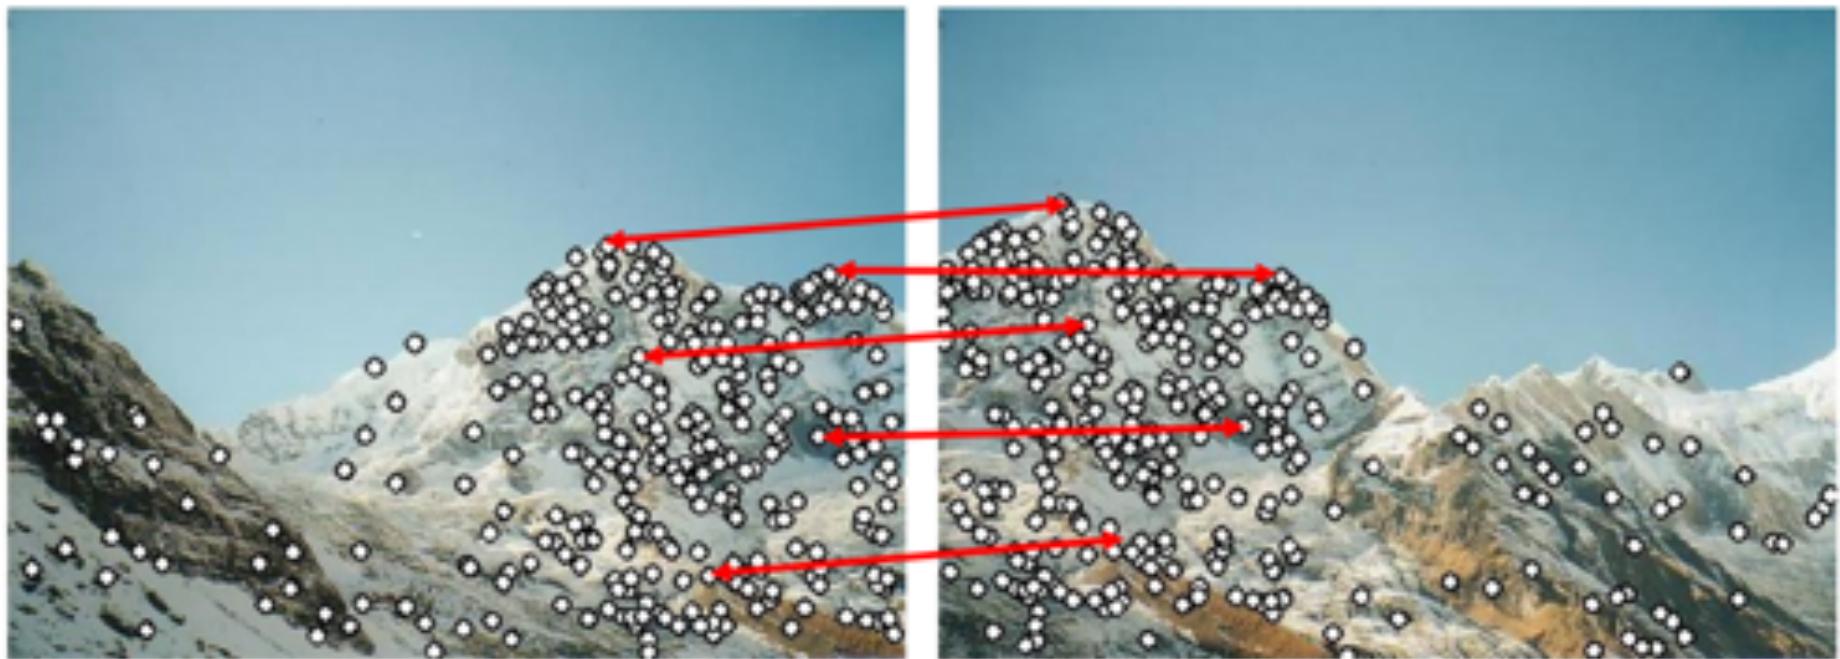
Applications

- Features are used for:
 - Image alignment (e.g., panoramic mosaics)
 - Object recognition
 - 3D reconstruction (e.g., stereo)
 - Motion tracking
 - Indexing and content-based retrieval
 - Robot navigation
 - ...

What is a keypoint?

What is a good feature?

- need to detect that feature in the next frame again



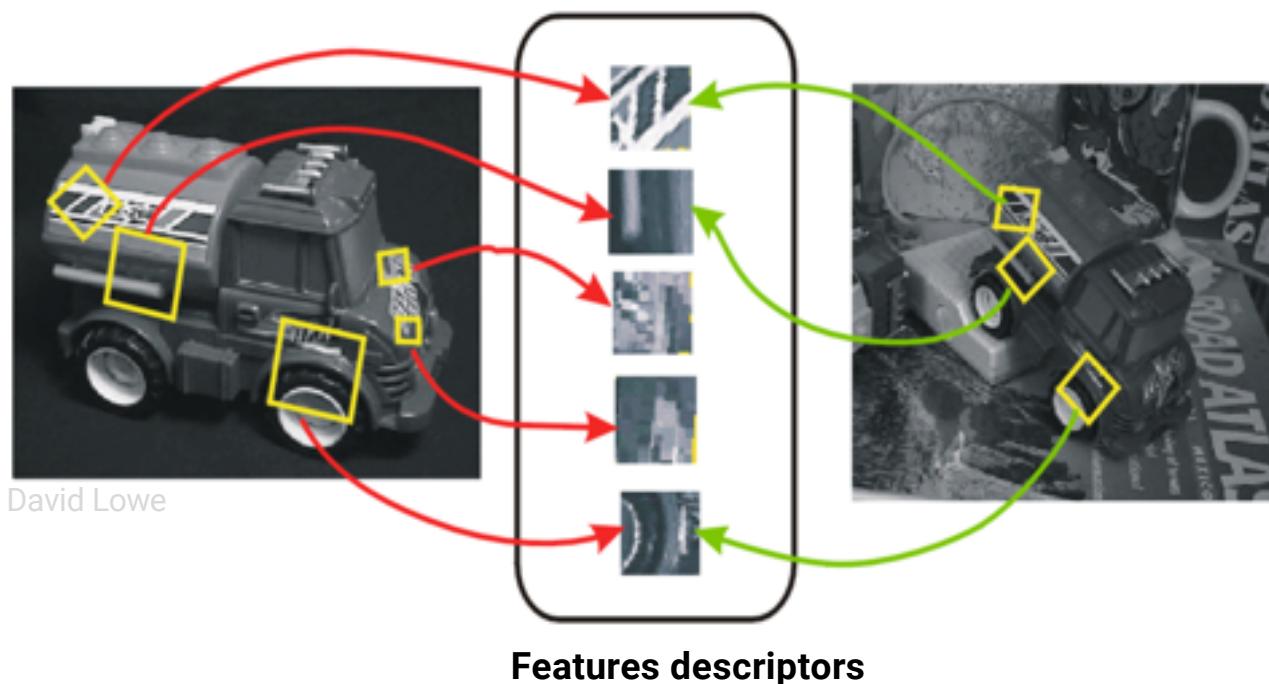
David Lowe

What makes a good feature?

- Want uniqueness
 - Leads to unambiguous matches in other images
- Look for “interest points”
- Stable under lighting and viewpoint changes
- Keypoints that are similarity- and affine-invariant, or at least stable
- Enable accurate matching of keypoints between images
 - Object recognition / Marker detection / image registration

Transformation invariance

- Want features that are invariant to transformations (affine-invariant)
 - geometric invariance: translation, rotation, scale
 - photometric invariance: brightness, exposure, ...



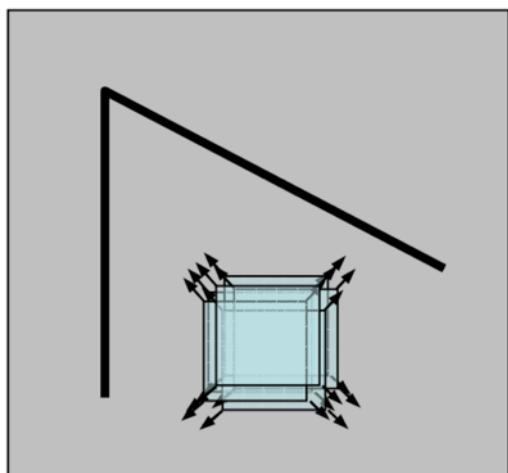
Feature Detection / Location

- What could be a possible detector?
 - random?
 - uniform?
 - edges?
 - **corners?**
- corners – intersection of edges
 - point where direction of intersections change
 - gradient of the image (in both directions) has high variation
 - this can be detected = we look for variation

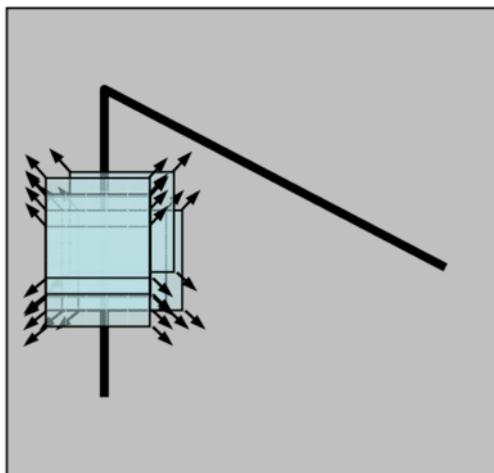
Harris corner detector – The basic idea

- C.Harris, M.Stephens. “A Combined Corner and Edge Detector”. 1988
- We should easily recognize the point by looking through a small window
- Shifting a window in any direction should give a large change in intensity

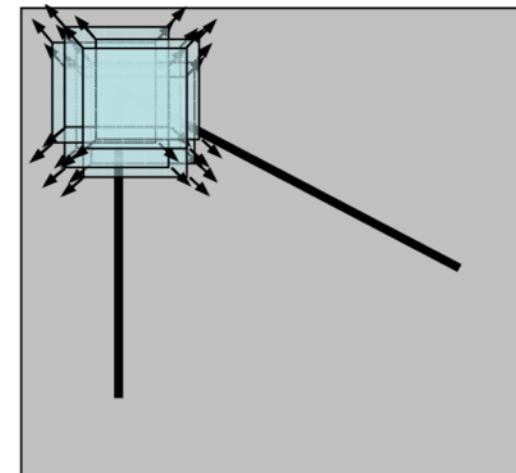
<https://courses.cs.washington.edu/>



“flat” region:
no change in all directions



“edge”: no change
along the edge direction



“corner”: significant
change in all directions

Harris corner detector Demo

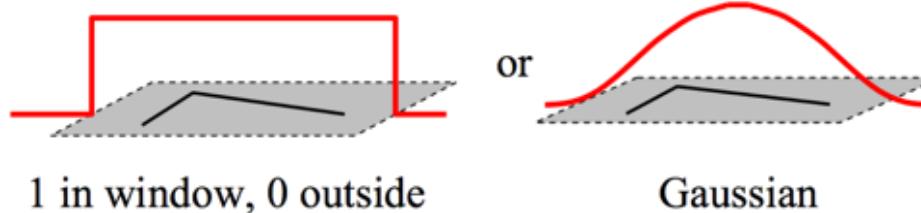
Harris corner detector – The math

- detect change of intensity for a window shift

$$E(u, v) = \sum_{x,y}^w [I(x + u, y + v) - I(x, y)]^2$$

↓
Window function ↓
Shifted intensity ↓
Intensity

- Window function $w(x,y)$:



<https://courses.cs.washington.edu/>

- Maximize variation within window

Harris Corner Detector – the Math

First order approximation from Taylor Series

$$f(x + u, y + v) \approx f(x, y) + u f_x(x, y) + v f_y(x, y)$$

$$\begin{aligned} E(u, v) &= \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{x,y} w(x, y) [I(x, y) + u I_x(x, y) + v I_y(x, y) - I(x, y)]^2 \\ &= \sum_{x,y} w(x, y) u^2 I_x^2 + 2uv I_x I_y + v^2 I_y^2 \\ &= \sum_{x,y} (u \quad v) \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \end{aligned}$$

rewrite as matrix equation

just the product of components of the gradient I_x, I_y
also known as structure tensor

Harris corner detector – The math

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

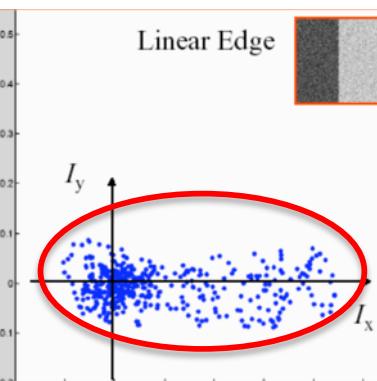
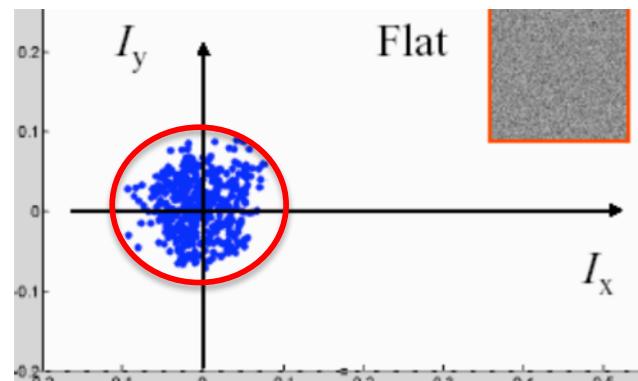
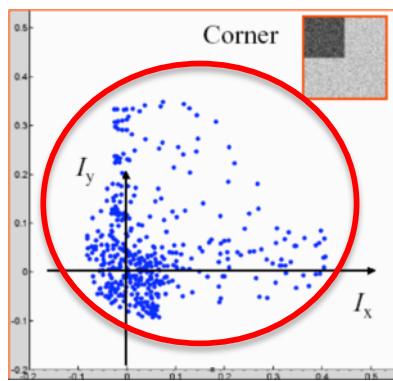
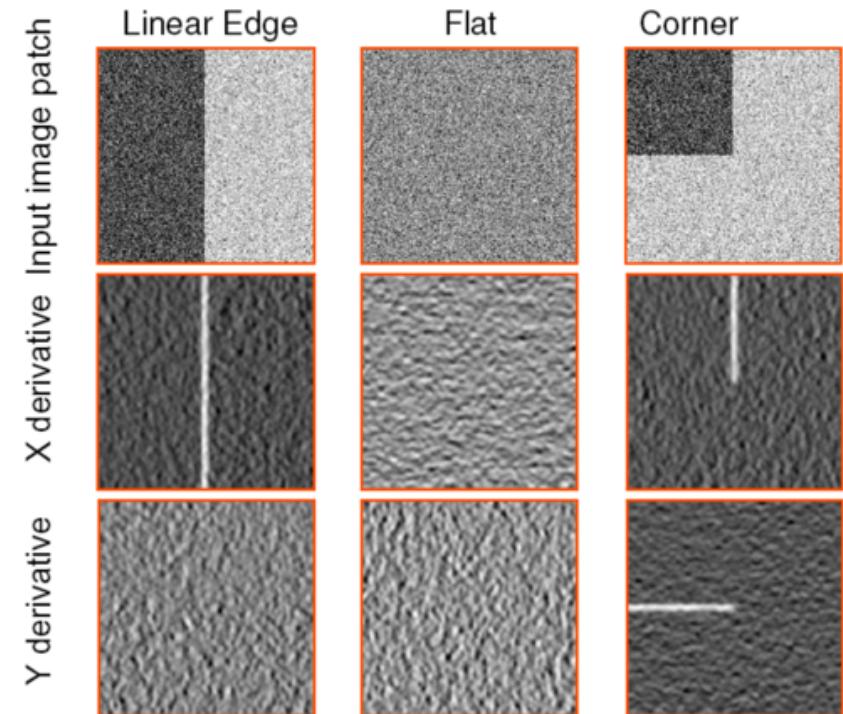
$$E(u, v) = \begin{pmatrix} u & v \end{pmatrix} M \begin{pmatrix} u \\ v \end{pmatrix}$$

- where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix}$$

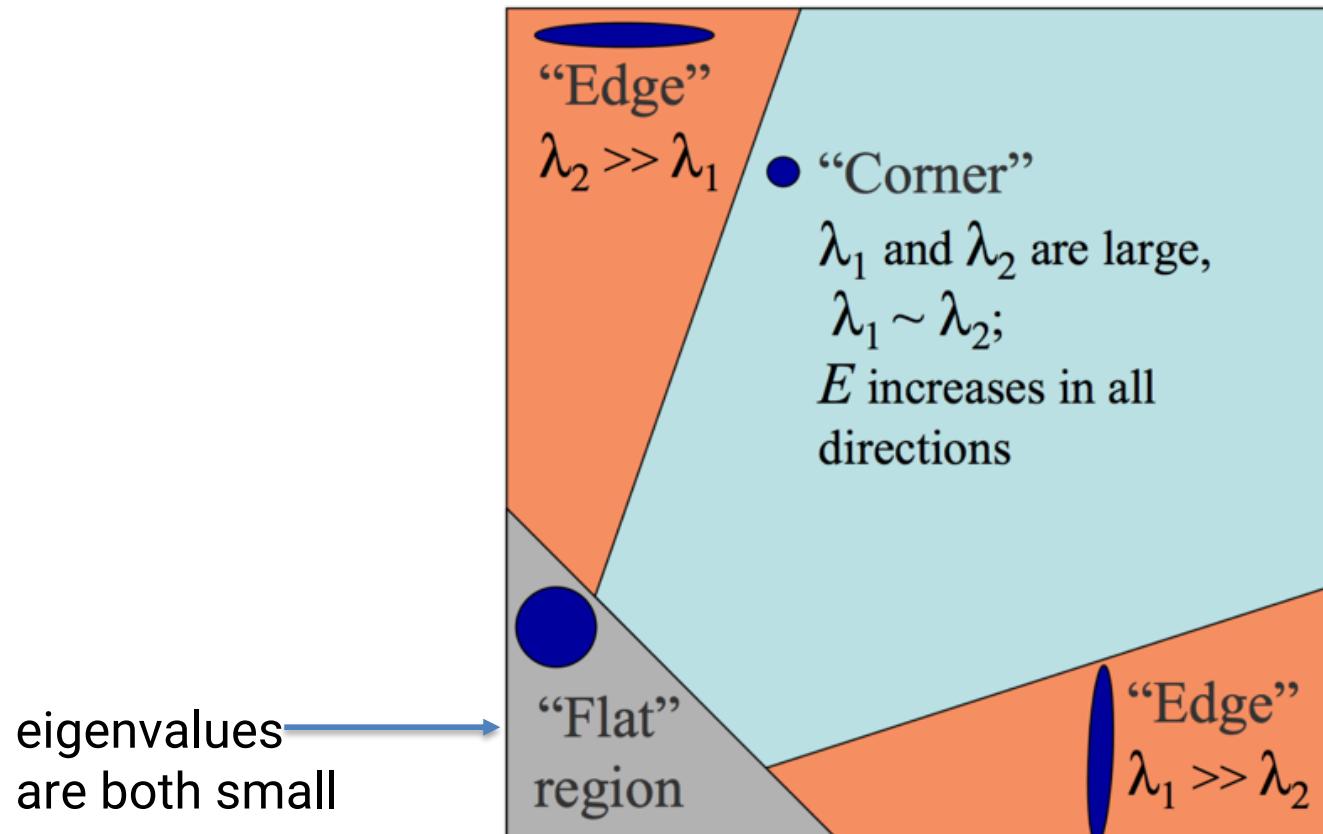
Harris corner detector – The intuitive way

- Treat gradient vectors as a set of (dx,dy) points with a center of mass defined as being at $(0,0)$
- Fit an ellipse to that set of points via scatter matrix
- Analyze ellipse parameters for varying cases



Harris corner detector

- Classification of image points using eigenvalues of M ???



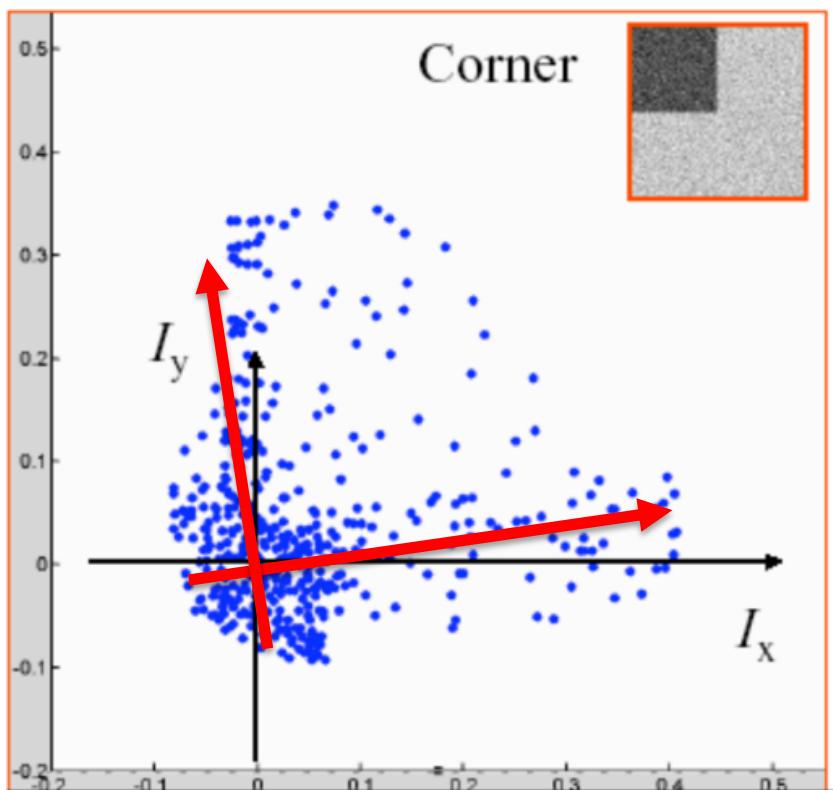
<https://courses.cs.washington.edu/>

EigenWhat?

Eigenvectors / Eigenvalues

Eigenvectors - Definition

$$A\mathbf{v} = \lambda\mathbf{v}$$

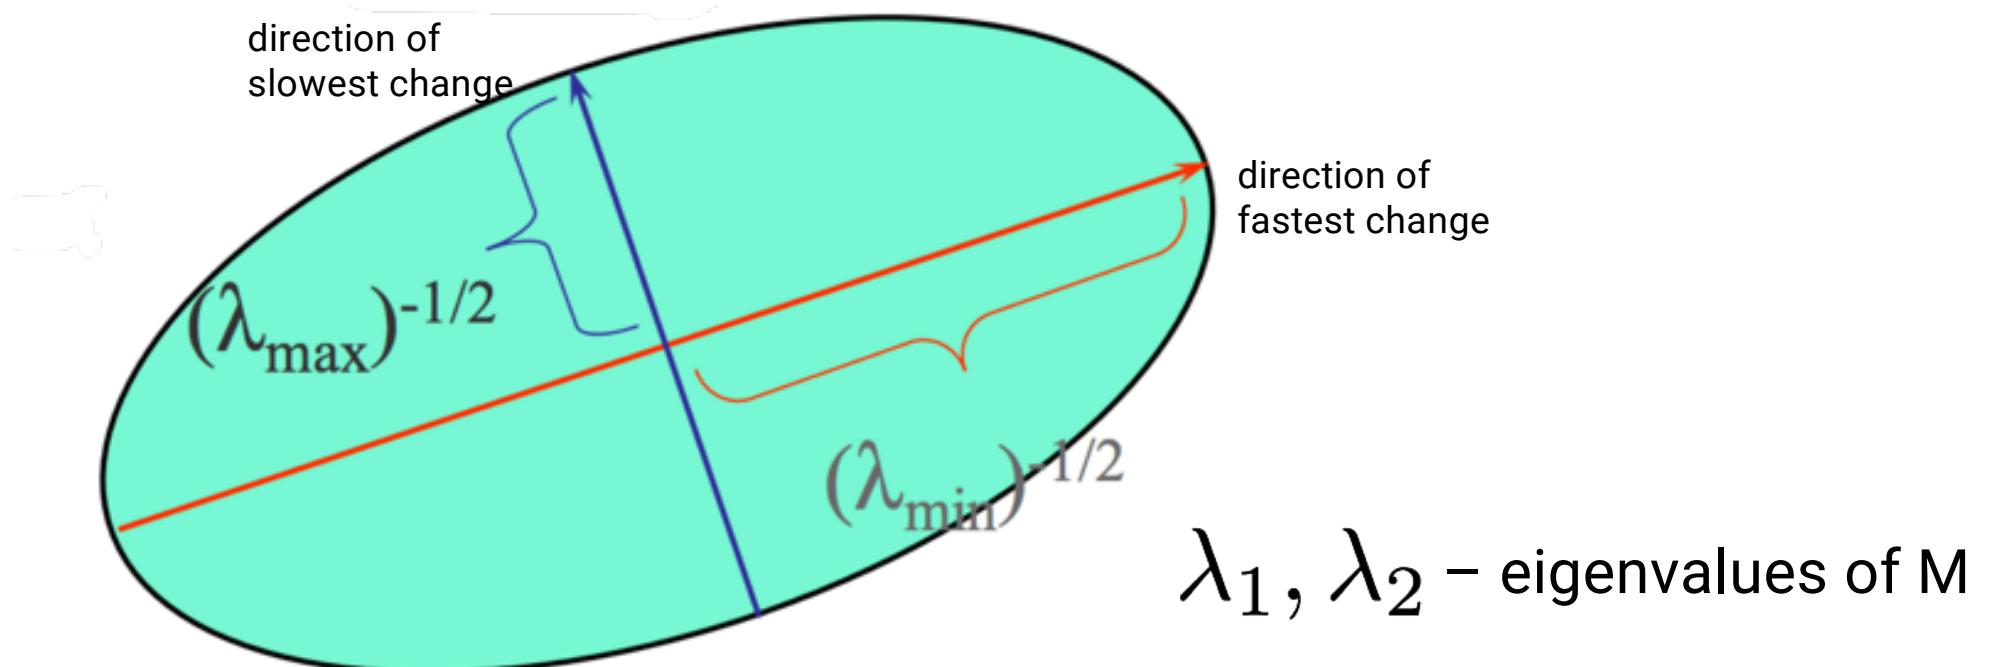


Harris corner detector – The math

- Intensity change in shifting window: eigenvalue analysis:

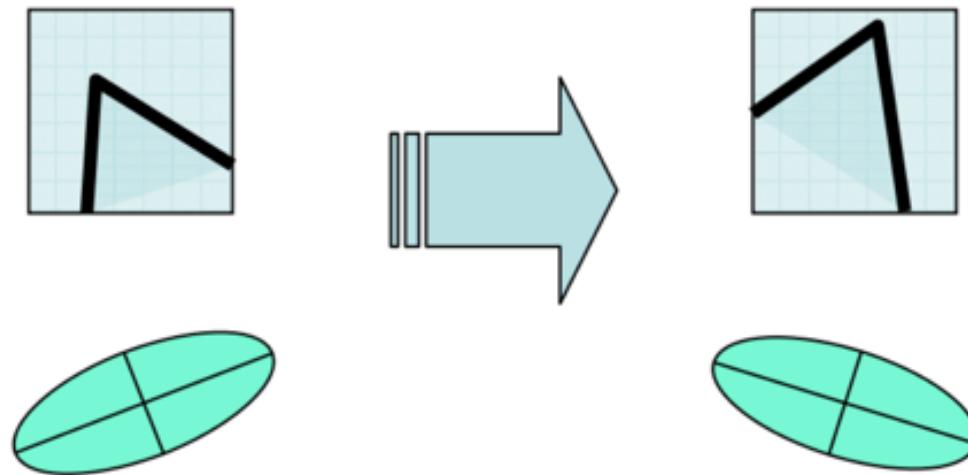
$$E(x, y) \simeq \begin{pmatrix} u & v \end{pmatrix} M \begin{pmatrix} u \\ v \end{pmatrix}$$

- Ellipse $E(u,v) = \text{const}$



Harris corner detector – Rotation invariance

- Ellipse rotates but its shape (i.e. eigenvalues) remains the same



<https://courses.cs.washington.edu/>

Harris corner detector – Efficient implementation

- One way to compute the eigenvalues efficiently

$$\lambda_{min} \approx \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2} = \frac{det(M)}{trace(M)}$$

- To avoid computing the eigenvalues which is computationally expensive one uses an approximation to determine if point is corner:

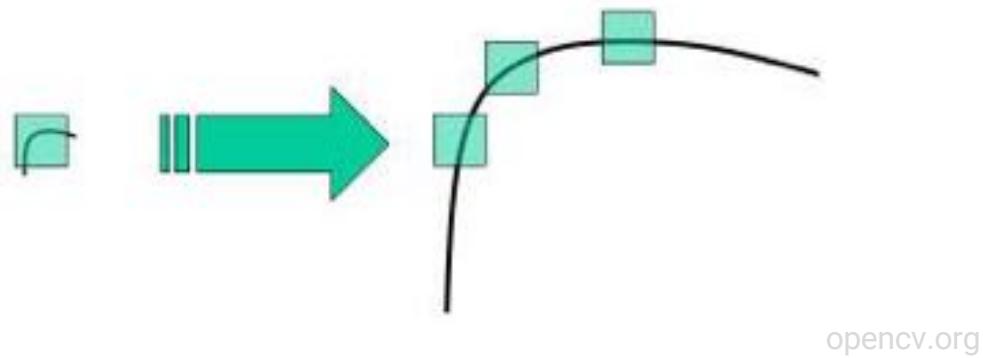
$$R = det(M) - k(trace(M))^2$$

Implementation is done as homework

Scale invariant feature transform (SIFT)

Scale invariant feature transform (SIFT)

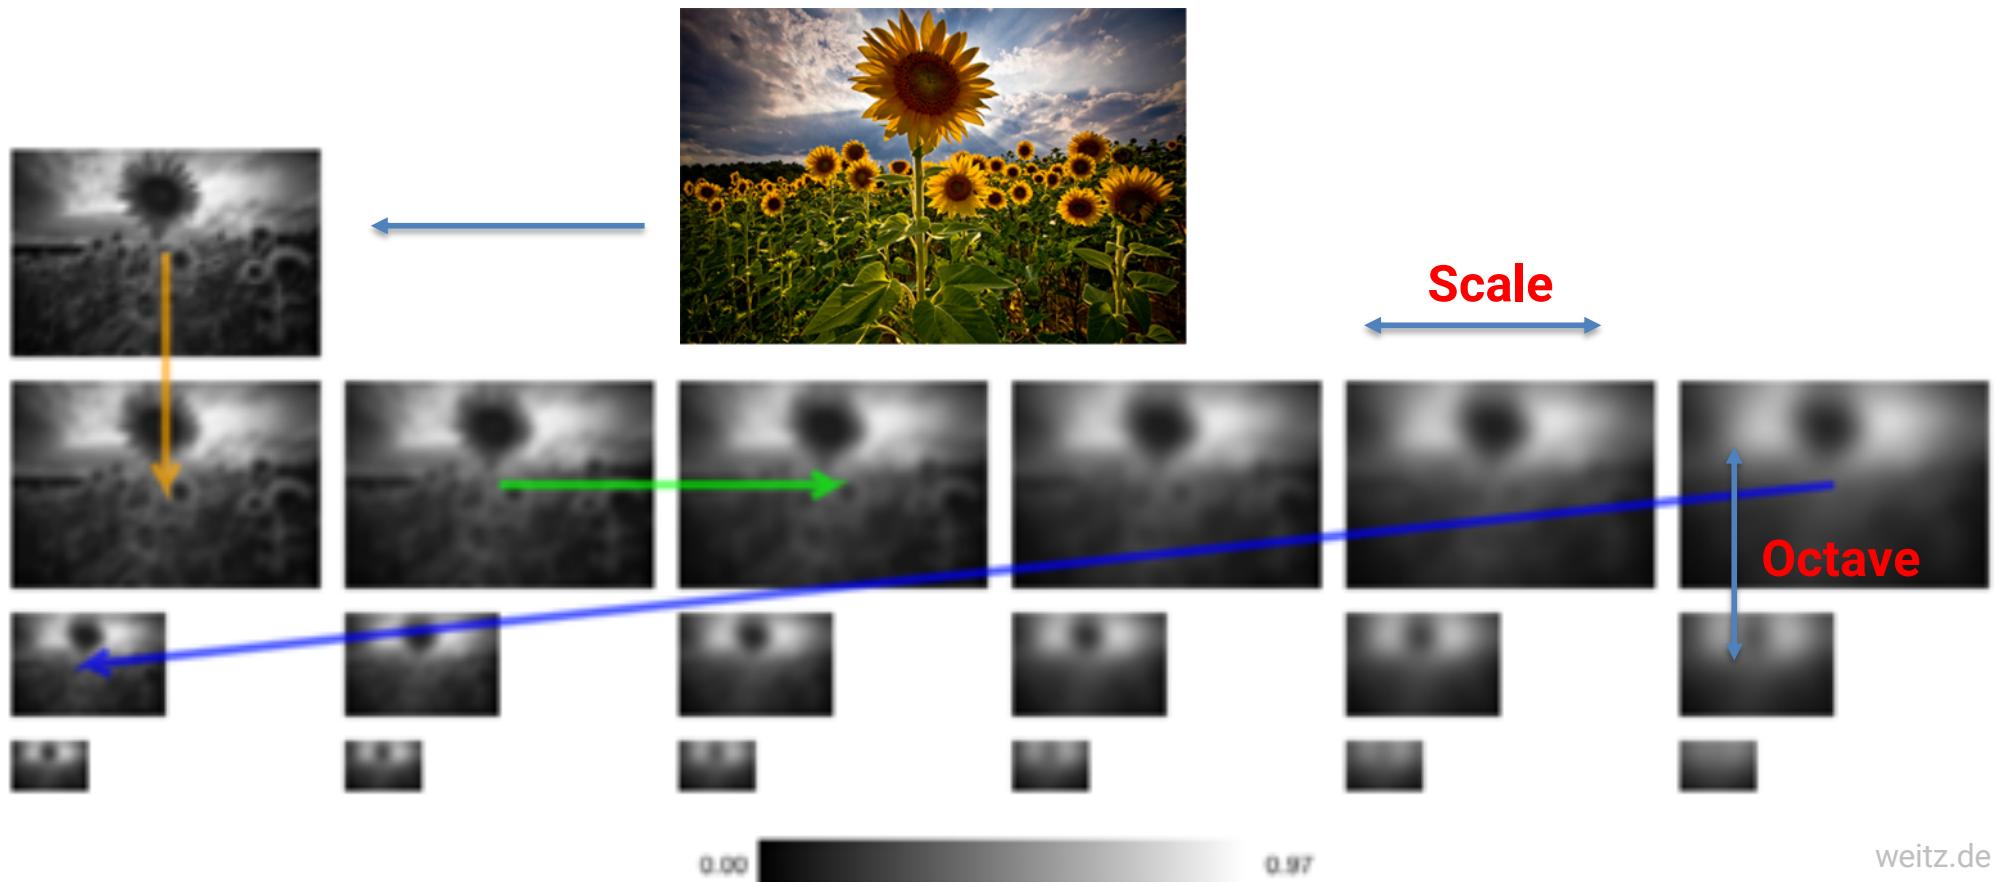
- Distinctive Image Features from Scale-Invariant Keypoints, David Lowe, 2004
- corner detectors are rotation-invariant
 - even if the image is rotated, one can find the same corners
 - corners remain corners in rotated image
- this is not true for scaling



opencv.org

- Harris detector is not scale-invariant

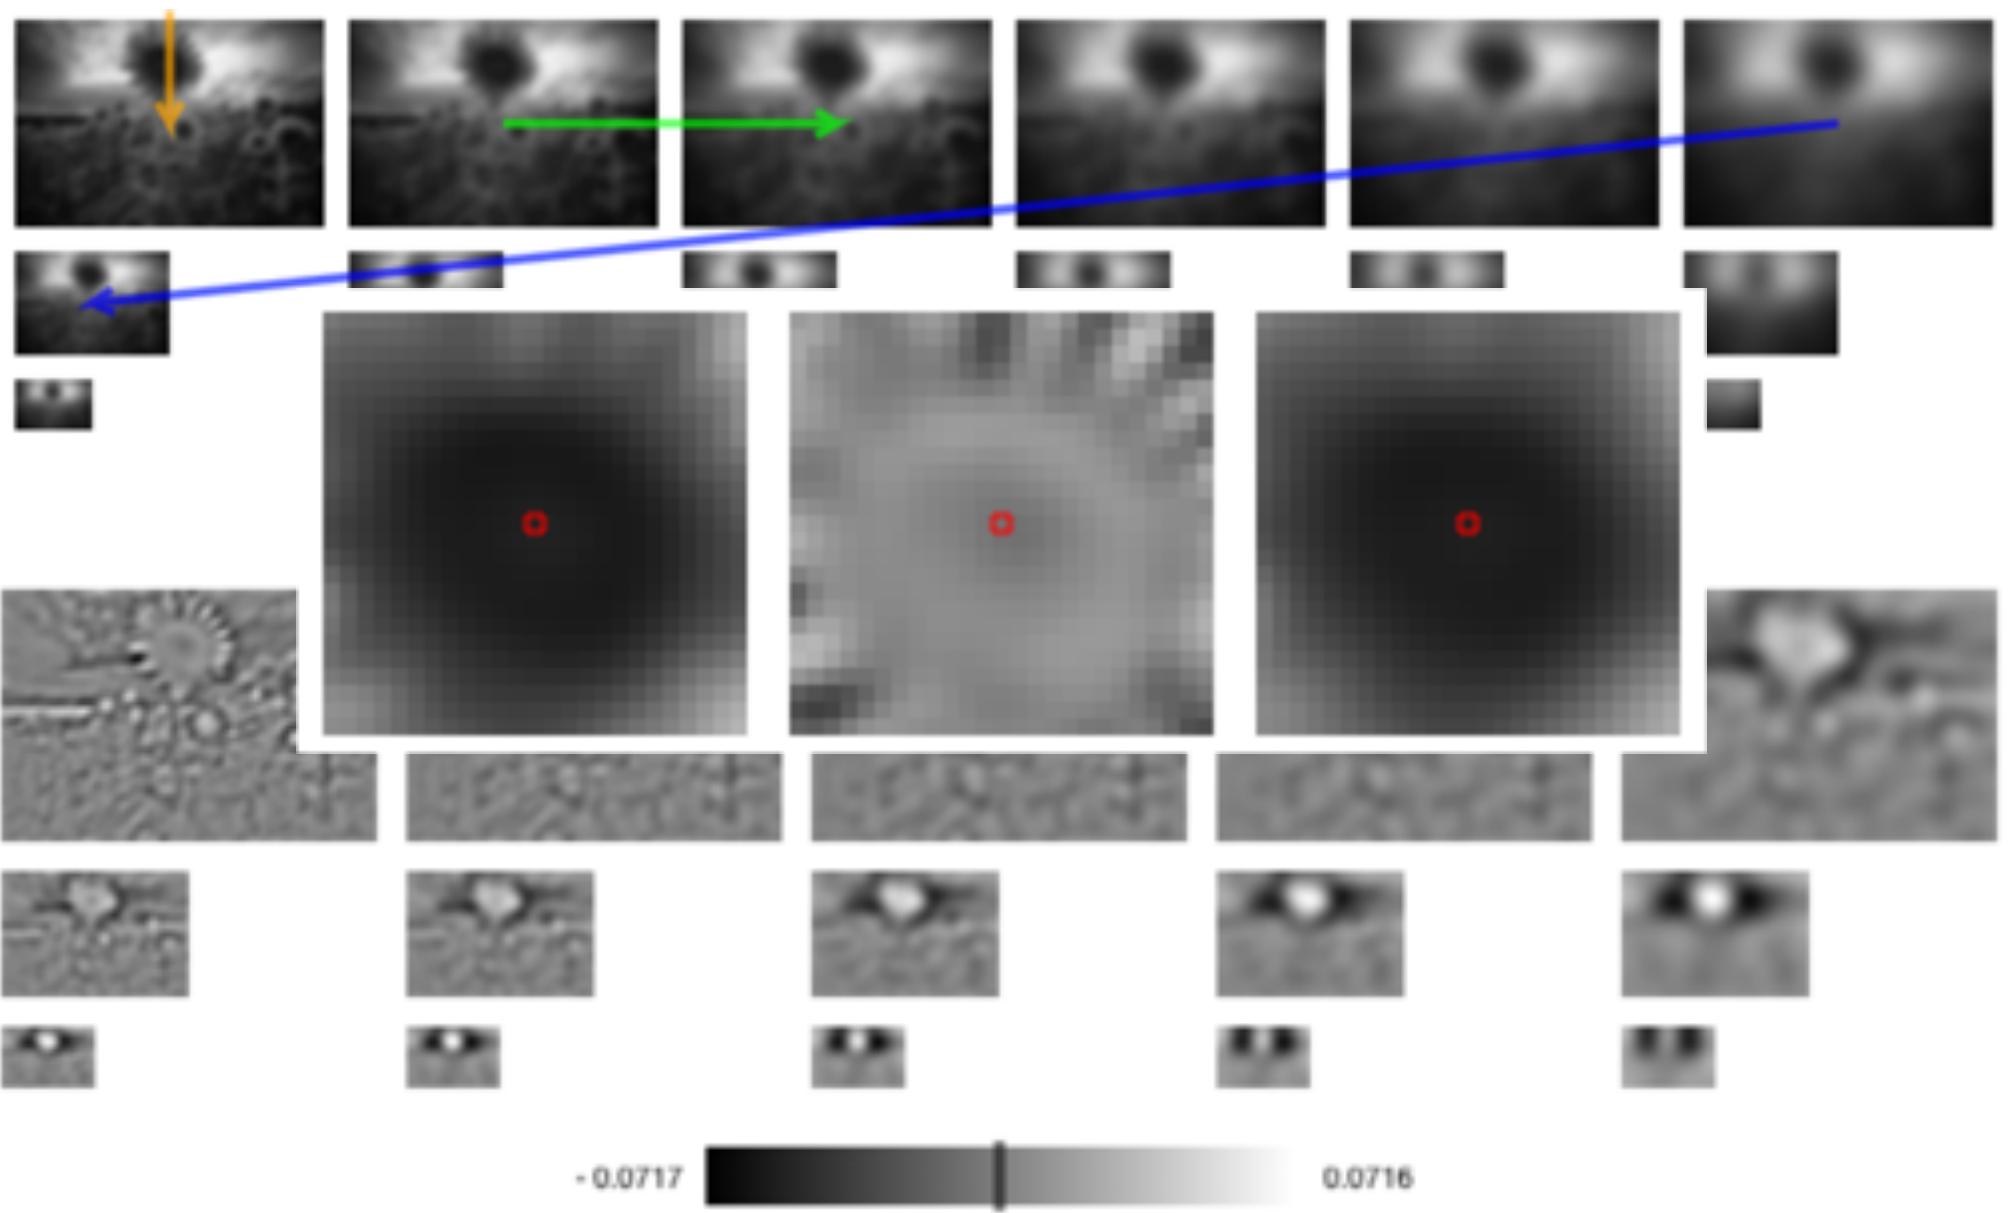
SIFT - Scale-space Extrema Detection



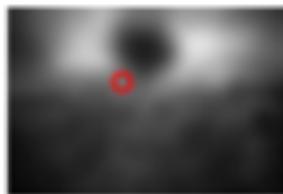
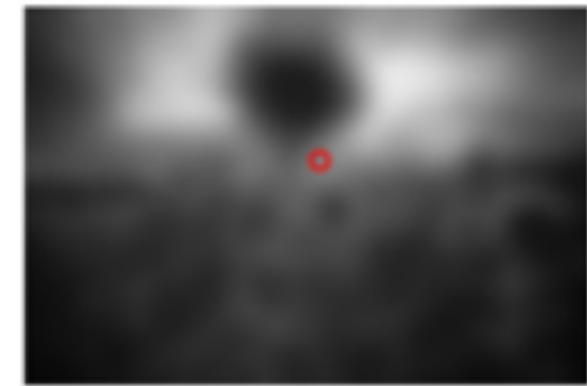
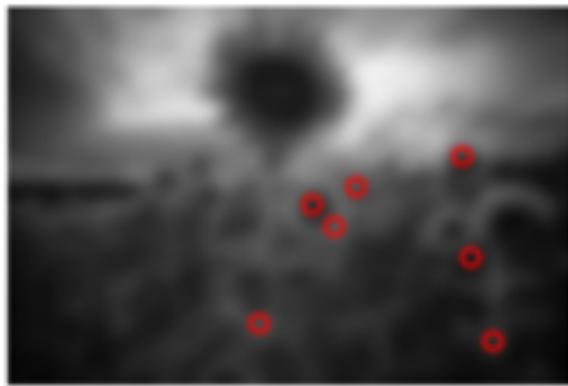
**Scale-space using bilinear interpolation
and gaussian convolution**

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$

SIFT - Difference of Gaussians



SIFT - Keypoints localization



weitz.de

- one pixel in an image is compared with its 8 neighbours as well as 9 pixels in next scale and 9 pixels in previous scales
- If it is a local extrema, it is a potential keypoint
- It basically means that keypoint is best represented in that scale

**Feature
Description of the image
window around the keypoint**

SIFT - Keypoint Descriptor

- 16x16 neighbourhood around the keypoint
- divided into 16 sub-blocks of 4x4 size
- each sub-block, 8 bin orientation histogram is created
 - total of 128 bin values are available
- represented as a vector to form keypoint descriptor

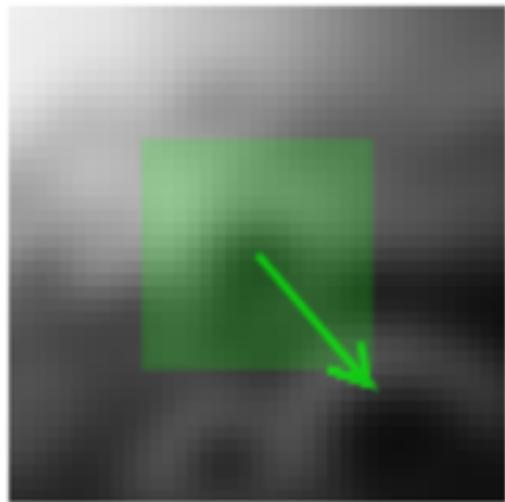
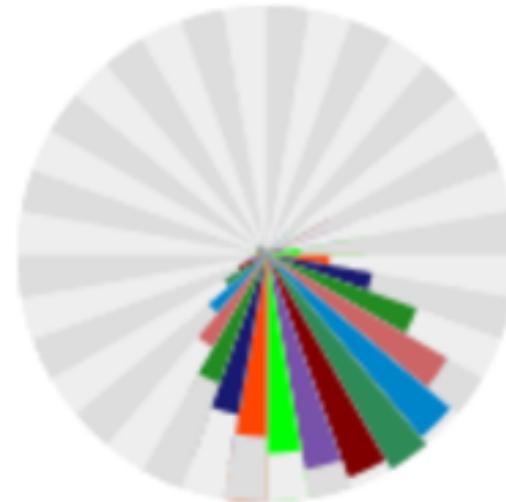


Image Gradients



weitz.de

Keypoint descriptor

SIFT keypoint detector Demo

Intermediate take away

Image feature example applications

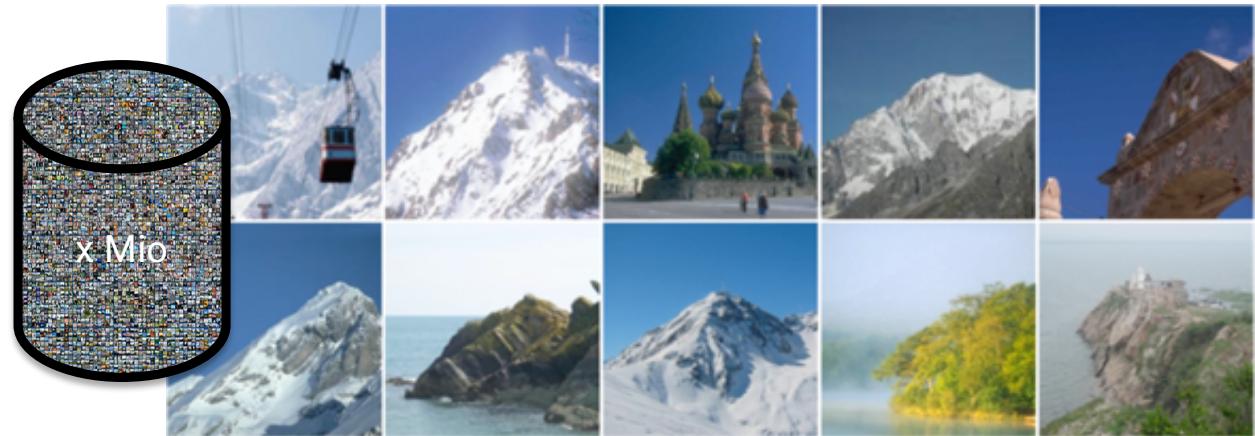
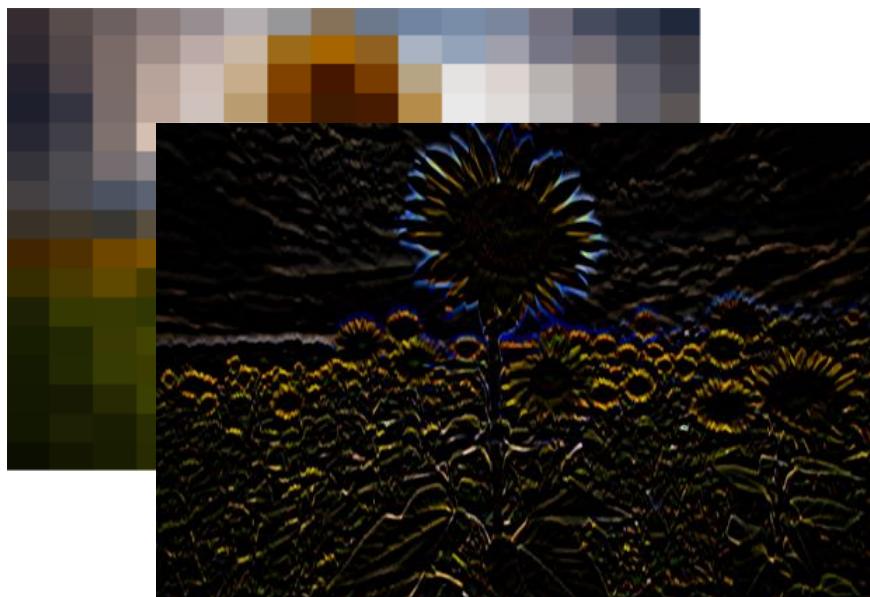


Image features



Global Descriptor

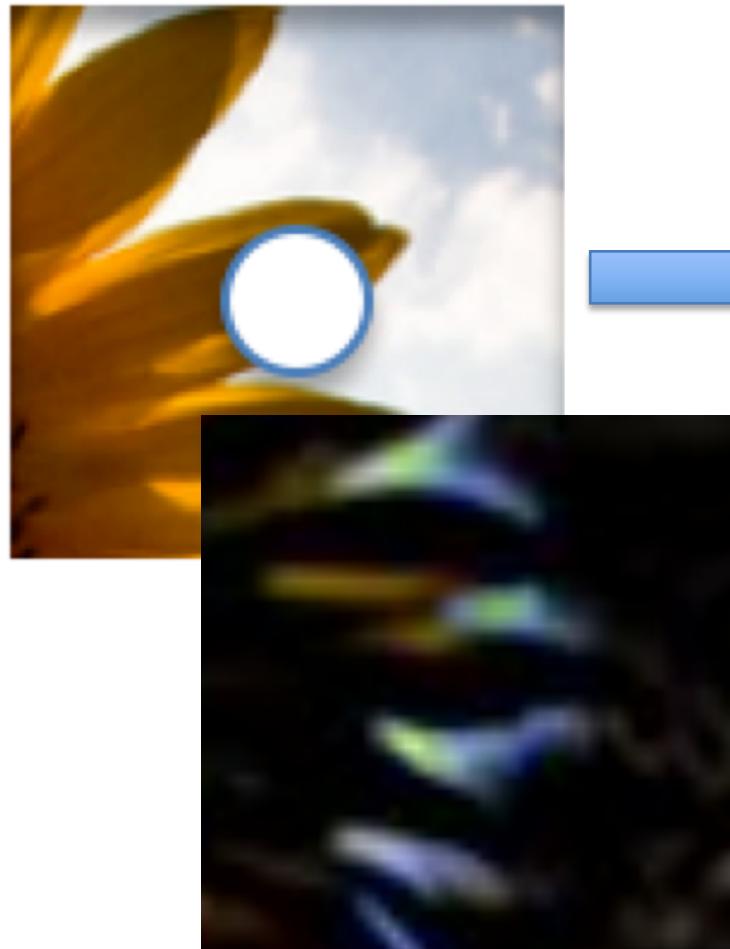


Local Descriptor

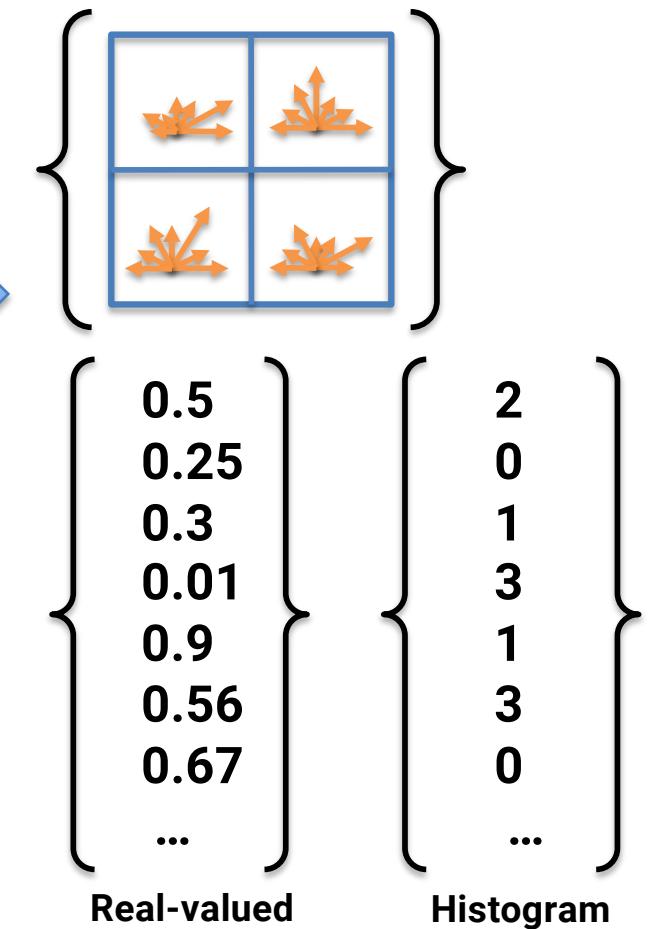


Image as a vector

What is the advantage of local feature information?



Feature
description

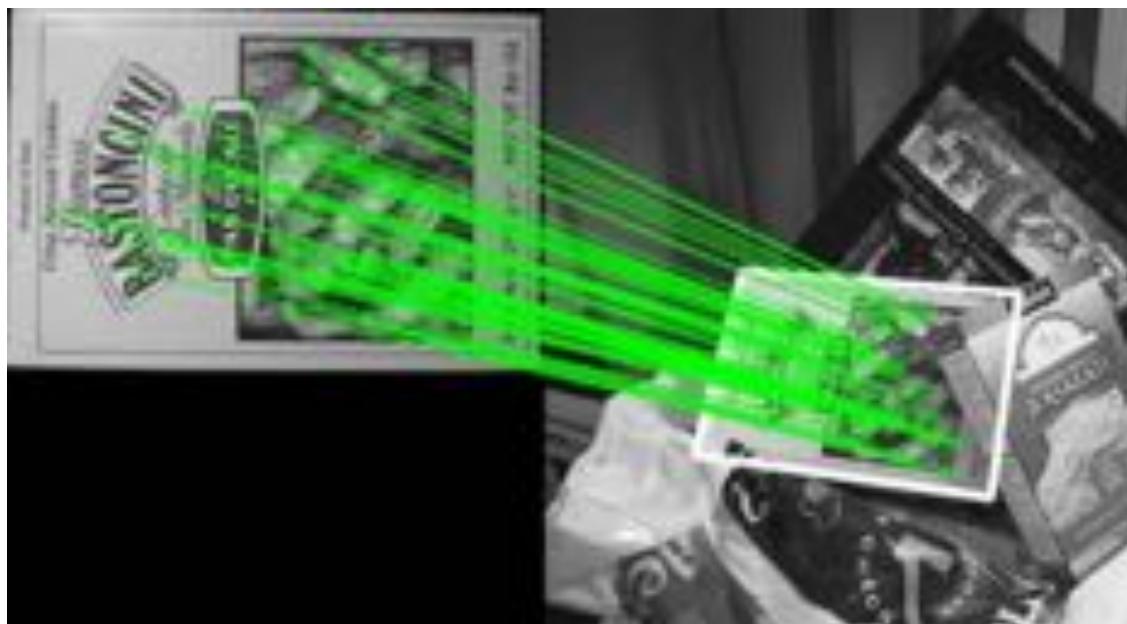


Real-valued

Histogram

Feature Matching

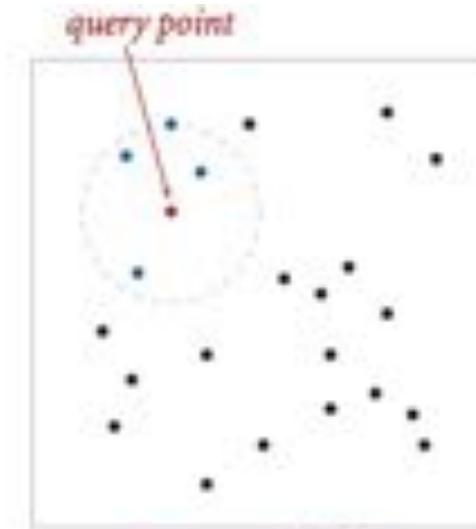
Feature Matching



<http://www.ubc.ca>



set of points



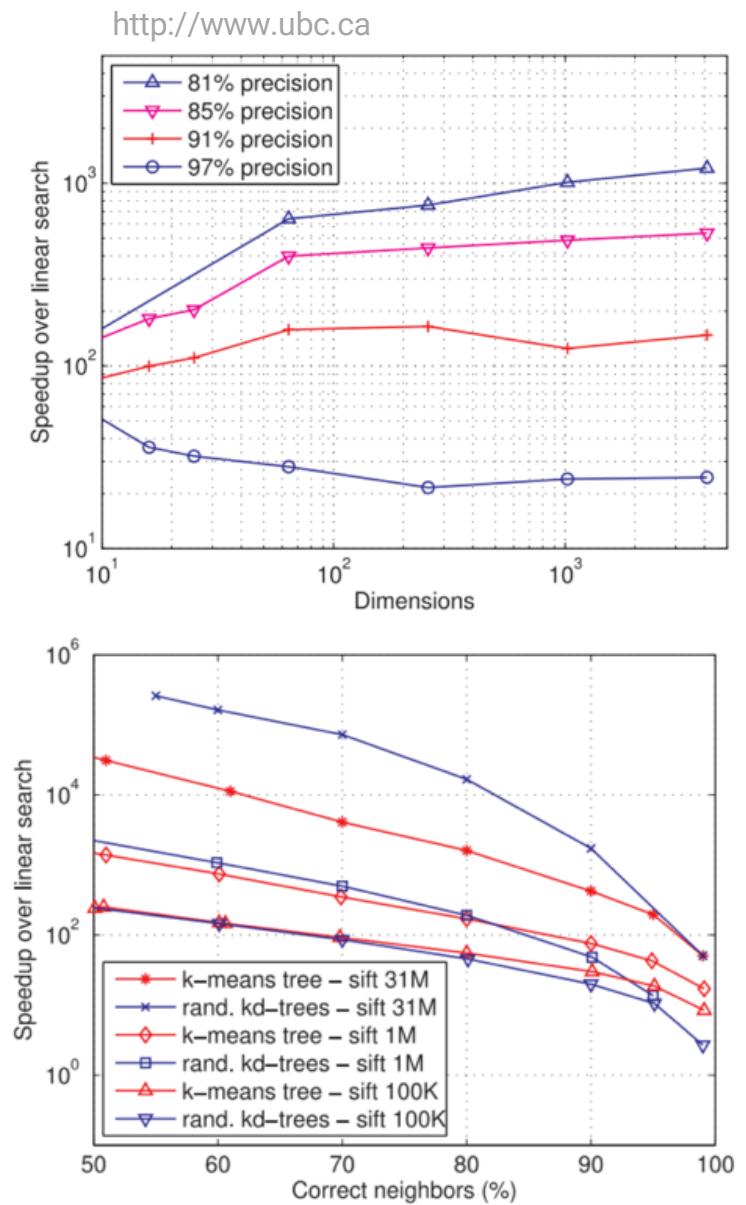
4 nearest neighbors

Feature Matching

- searching for the **most similar matches** to high-dimensional vectors  **nearest neighbor matching**
- **no known exact algorithms** for solving these high-dimensional problems that are faster than linear search  **Curse of dimensionality**
- **efficient algorithm** for fast nearest neighbor (NN) matching in large data **sets can bring speed improvements of several orders of magnitude**
- **FLANN – Fast Library for Approximate Nearest Neighbor**
 - library for performing fast approximate nearest neighbor searches in high dimensional spaces

Feature Matching

- Obtain **speed improvement by allowing an approximate search**
 - not all the neighbors returned are exact
 - some are approximate but still close to the exact neighbors
- **Approximate nearest neighbor** search algorithms provide **~95% correct neighbors**
 - **BUT:** are two or more orders of magnitude faster than linear search



Feature Matching Applications

- Nearest neighbor search problem is also of major importance for many **applications**:
 - machine learning
 - image- / document-retrieval
 - data compression
 - bio-informatics
 - data analysis
 - computer vision

Demo – Image stitching

Feature Matching

- Given a set of points in metric space M and a query point q

$$P = \{p_1, p_2, p_3, \dots, p_n\} \quad q \in M$$

- Find elements that is closest to q with respect to a metric distance d

$$NN(q, P) = \operatorname{argmin}_{x \in P} d(q, x)$$

Feature Matching

- **How to compare feature vectors in high-dimensional space?**
 - Distance Metric
- How to compare feature vectors in high-dimensional space **efficiently**?
 - Spatial Datastructure
 - Approximate Nearest Neighbor

Distance Metrics

Distance metric

- **Euclidean distance**

- ‘straight-line’ distance between points (L^2 -norm)

$$d_{l2}(p, q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

- **Manhattan distance**

- L1-norm

$$d_{l1}(p, q) = \sum_{i=1}^n |p_i - q_i|$$

- **Maximum norm**

- L^∞

$$d_{l\infty}(p, q) = \max_i(|p_i - q_i|)$$

Distance metric

- **Hamming distance**
 - distance between binary code
 - "foobar" and "f00bat" is 3.
 - 1011101 and 1001001 is 2.
 - 2173896 and 2233796 is 3.
- **Cosine similarity (no proper distance metric)**
 - similarity between two vectors (inner product, $\cos(\text{angle})$ between vectors)
 - distance between histograms

$$a \cdot b = \|a\| \|b\| \cos \theta$$

Fast Approximate Nearest Neighbor

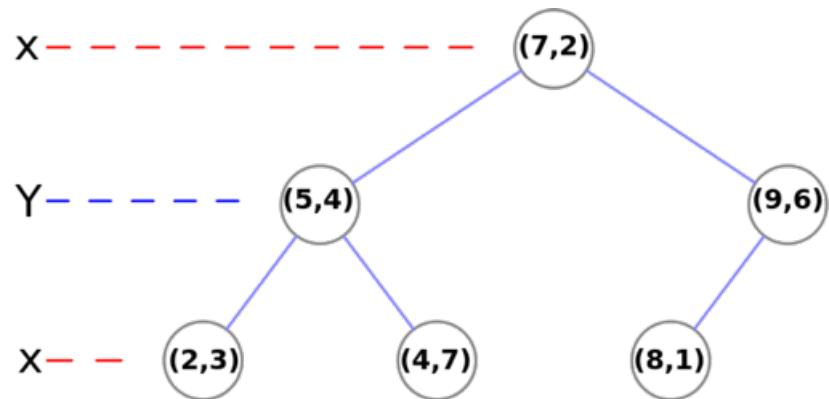
Fast Approximate Nearest Neighbor (FLANN)

- Library for performing fast **approximate nearest neighbor** searches in **high dimensional spaces**
 - <http://www.cs.ubc.ca/research/flann/>
 - collection of standard algorithms that work best in practice
 - written in c++ with binding for Matlab and Python
- Algorithm evaluation showed two best performing algorithms:
 - **Multiple randomized k-d trees**
 - **Priority search k-means tree**

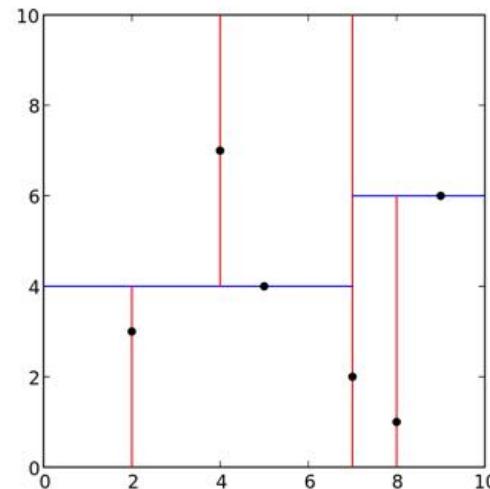
Multiple Randomized kd-Trees

kd-tree (Recap)

- **spatial data structure** for organizing points in **high dimensionality spaces**
- **binary tree** where each node is a k-d point
- every **node** can be seen as a **plane dividing space in two parts**
- splitting along changing dimensions / axis
- **results in balanced k-d tree**

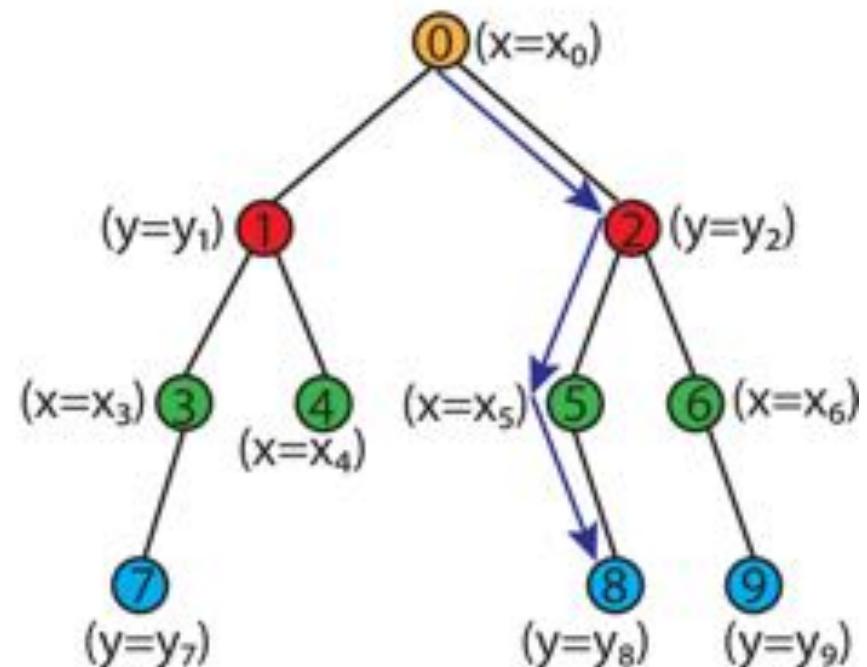
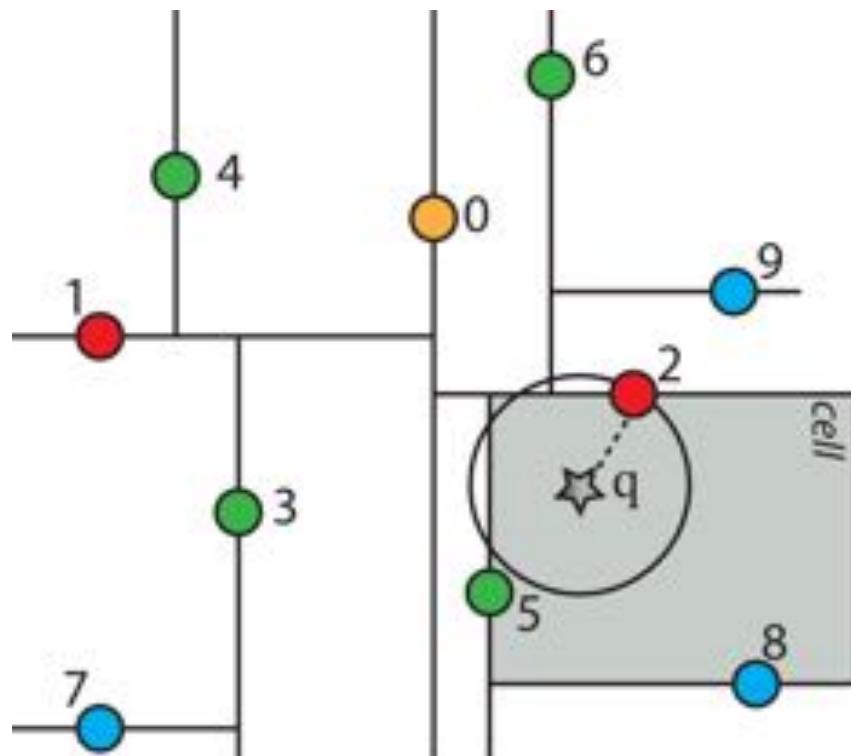


white board example



<http://wikipedia.org>

Query in kd-trees



kd-tree

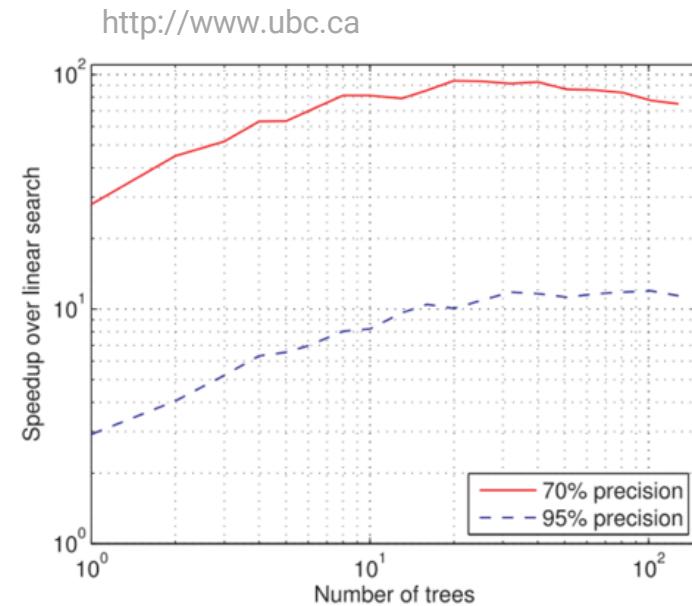
- Very **effective in low dimensional spaces**
- Performance **decreases for high dimensional data**
 - requires searching a large number of nodes
- Speedup by approximating queries:
 - ‘**error bound**’ approximate search
 - ‘**time bound**’ approximate search

Randomized kd-Trees

- **Idea:** multiple randomized k-d trees which are searched in parallel
 - build scheme similar to classic kd-tree
- **Difference to classic kd-tree:**
 - classic – split dimensions are chosen by highest variance
 - randomized – split dimensions are chosen randomly from the **top 5 dimensions highest variance**
- In practice often **20 trees** used

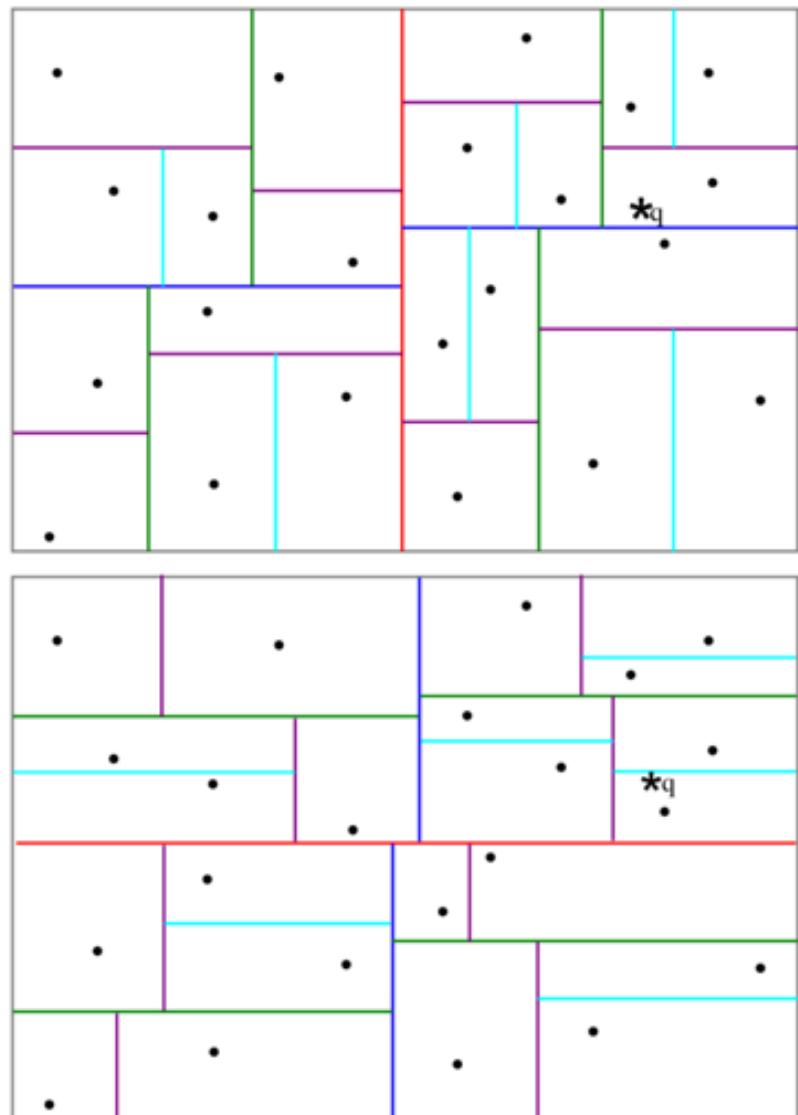
Searching randomized k-d forest

- Maintain **single priority queue across forest**
- PQ is ordered by increasing distance
 - the search will explore first the closest leaves from all the trees
- examined **points are marked as visited across all trees** (not visited in another tree again)
- approximation is determined by **maximum number of visited leaf nodes**



100k SIFT feature data set

Randomized kd-Trees



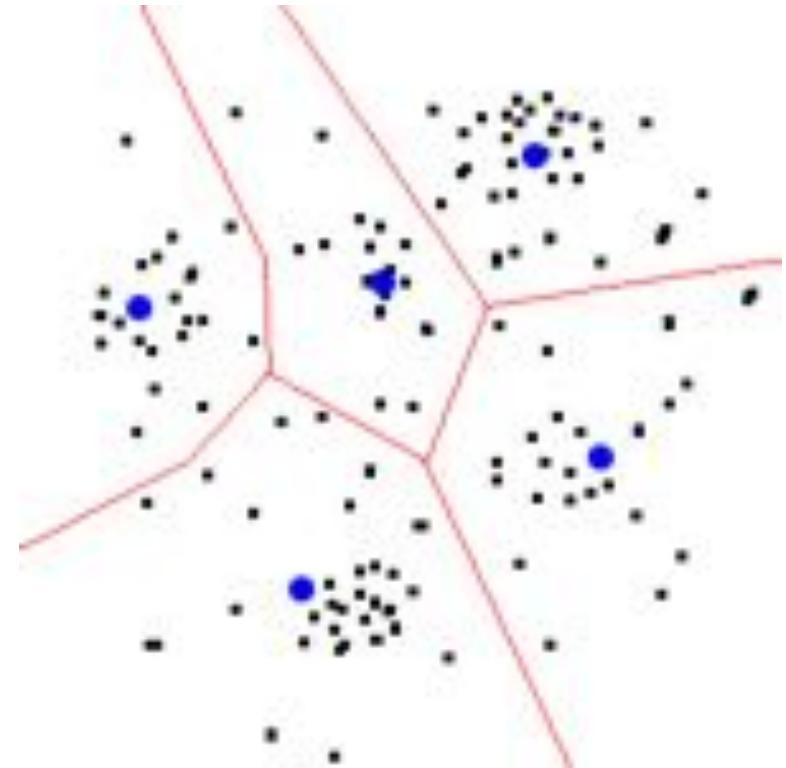
Priority Search k-means Tree

k-means

- important **clustering method** in CS
- partitions **N data points** in **k clusters**
 $S = \{S_1, S_2, \dots S_k\}$
- each **data point belongs to cluster with nearest mean**
 - mean of the cluster is cluster center/representative

$$\arg \min_s \sum_{i=1}^k \sum_{x \in s_i} \|x - \mu_i\|^2$$

- **results in partitioning of data space (voronoi cells)**
- also known as **Lloyd's Algorithm**



<http://mnemstudio.org/>

k-means Trees

- **Construction:**
 - Tree is constructed by **partitioning the data points at each level into K distinct regions** using k-means clustering
 - Apply recursively to the points in each region
 - Recursion is stopped when the number of points in a region is smaller than K
- **Search:**
 - initially traversing the tree from the root to the closest leaf (following closest cluster centre)
 - add unvisited branches to priority queue
 - use branch that has the closest center to the query point
 - restart the tree traversal from that branch
 - stop after number of visited leaf nodes

k-means Trees

