

Machine Learning

Lecture 3

Simple Classifiers: Nearest Centroids and Linear Classification

Felix Bießmann

Beuth University & Einstein Center for Digital Future




Overview of today's lecture

- Today we will introduce three simple classifiers
 1. Nearest Centroid Classifier (NCC)
 2. Perceptron
 3. K-Nearest Neighbor (KNN)
- These algorithms are extremely powerful
- Often they can compete with complex algorithms
- Some aspects can be motivated by biological cognition



Prototypes: Psychological Models of Abstract Ideas



Psychologists postulated that we learn **prototypes** [??]


Toy data example:

Two dimensional input $\mathbf{x} \in \mathbb{R}^2$

Two *classes* of data, Δ and \circ



Prototypes: Psychological Models of Abstract Ideas



Prototypes μ_{Δ} and μ_o can be the class means

$$\mu_{\Delta} = 1/N_{\Delta} \sum_n^{N_{\Delta}} \mathbf{x}_{\Delta,n}$$


$$\mu_o = 1/N_o \sum_n^{N_o} \mathbf{x}_{o,n}$$

Distance from w_{Δ} to new data \mathbf{x}

$$\|\mu_{\Delta} - \mathbf{x}\|_2$$



Prototypes: Psychological Models of Abstract Ideas



For new data x check:
Is x more similar to μ_o ?

$$\|\mu_{\Delta} - x\| > \|\mu_o - x\|$$

yes? $\rightarrow x$ belongs to μ_o

no? $\rightarrow x$ belongs to μ_{Δ}

This is called a
nearest centroid classifier



Nearest Centroid Classification Algorithm (Batch Mode)

Algorithm 1 Computation of Class-Centroids

Require: data $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^D$, labels $y_1, \dots, y_N \in \{1, \dots, K\}$

Ensure: Class means $\boldsymbol{\mu}_k$, $k \in \{1, \dots, K\}$

- 1: # Initialize means and counters for each class
 - 2: # Computation of class means
 - 3: **for** Class $k = 1, \dots, K$ **do**
 - 4: $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{i=1}^{N_k} \mathbf{x}_i$
 - 5: **end for**
-



Batch Computations vs. Streaming

Solutions for algorithms can be obtained

- In Batch Mode:
 - Use all available data at once
 - Requires to store all data in memory
- In Streaming Mode:
 - Use one data point at a time
 - Requires to store **only** centroids



Iterative Computation of the Mean

Given the mean μ_{N-1} computed from $N - 1$ samples we want to update μ_{N-1} with the N th sample \mathbf{x}_N to obtain μ_N

$$\begin{aligned}\mu_N &= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n \\ &= \frac{1}{N} \sum_{n=1}^{N-1} \mathbf{x}_n + \frac{1}{N} \mathbf{x}_N \\ &= \frac{N-1}{N} \underbrace{\frac{1}{N-1} \sum_{n=1}^{N-1} \mathbf{x}_n}_{\mu_{N-1}} + \frac{1}{N} \mathbf{x}_N \\ &= \frac{N-1}{N} \mu_{N-1} + \frac{1}{N} \mathbf{x}_N\end{aligned}$$



Nearest Centroid Classification Algorithm (Streaming)

Algorithm 2 Iterative computation of Class-Centroids

Require: data $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^D$, labels $y_1, \dots, y_N \in \{1, \dots, K\}$

Ensure: Class means $\boldsymbol{\mu}_k$, $k \in \{1, \dots, K\}$

1: # Initialize means and counters for each class

2: $\forall k : \boldsymbol{\mu}_k = \mathbf{1} \cdot 0, N_k = 0$

3: # Iterative computation of class means

4: **for** Data point $i = 1, \dots, N$ **do**

5: # Update means and counters

6: $k = y_i$

7: $\boldsymbol{\mu}_k = \frac{N_k}{N_k+1} \boldsymbol{\mu}_k + \frac{1}{N_k+1} \mathbf{x}_i$

8: $N_k = N_k + 1$

9: **end for**



Nearest Centroid Classification

Algorithm 3 Nearest Centroid Prediction




Require: Data point $\mathbf{x} \in \mathbb{R}^D$, class centroids $\boldsymbol{\mu}_k$, $k \in \{1, \dots, K\}$

Ensure: Class membership k^*

- 1: # Compute nearest class centroid in discriminative subspace
 - 2: $k^* = \operatorname{argmin}_k \|\boldsymbol{\mu}_k - \mathbf{x}\|_2$.
-



Toy Data Example NCC



From Prototypes to Linear Classification

$$\begin{aligned} \text{distance}(\mathbf{x}, \mu_{\Delta}) &> \text{distance}(\mathbf{x}, \mu_o) \\ \|\mathbf{x} - \mu_{\Delta}\| &> \|\mathbf{x} - \mu_o\| \end{aligned} \tag{1}$$



From Prototypes to Linear Classification

$$\text{distance}(\mathbf{x}, \mu_{\Delta}) > \text{distance}(\mathbf{x}, \mu_o) \quad (1)$$

$$\|\mathbf{x} - \mu_{\Delta}\| > \|\mathbf{x} - \mu_o\|$$

$$\Leftrightarrow \|\mathbf{x} - \mu_{\Delta}\|^2 > \|\mathbf{x} - \mu_o\|^2$$

$$\Leftrightarrow \mathbf{x}^T \mathbf{x} - 2\mu_{\Delta}^T \mathbf{x} + \mu_{\Delta}^T \mu_{\Delta} > \mathbf{x}^T \mathbf{x} - 2\mu_o^T \mathbf{x} + \mu_o^T \mu_o$$

$$\Leftrightarrow \mu_{\Delta}^T \mathbf{x} - \mu_{\Delta}^2/2 < \mu_o^T \mathbf{x} - \mu_o^2/2$$

$$\Leftrightarrow 0 < \underbrace{(\mu_o - \mu_{\Delta})^T}_{\mathbf{w}} \mathbf{x} - 1/2 \underbrace{(\mu_o^T \mu_o - \mu_{\Delta}^T \mu_{\Delta})}_{\beta}$$



From Prototypes to Linear Classification

$$\text{distance}(\mathbf{x}, \mu_{\Delta}) > \text{distance}(\mathbf{x}, \mu_o) \quad (1)$$

$$\|\mathbf{x} - \mu_{\Delta}\| > \|\mathbf{x} - \mu_o\|$$

$$\Leftrightarrow \|\mathbf{x} - \mu_{\Delta}\|^2 > \|\mathbf{x} - \mu_o\|^2$$

$$\Leftrightarrow \mathbf{x}^T \mathbf{x} - 2\mu_{\Delta}^T \mathbf{x} + \mu_{\Delta}^T \mu_{\Delta} > \mathbf{x}^T \mathbf{x} - 2\mu_o^T \mathbf{x} + \mu_o^T \mu_o$$

$$\Leftrightarrow \mu_{\Delta}^T \mathbf{x} - \mu_{\Delta}^2/2 < \mu_o^T \mathbf{x} - \mu_o^2/2$$


$$\Leftrightarrow 0 < \underbrace{(\mu_o - \mu_{\Delta})^T \mathbf{x}}_{\mathbf{w}} - 1/2 \underbrace{(\mu_o^T \mu_o - \mu_{\Delta}^T \mu_{\Delta})}_{\beta}$$

Linear Classification

$$\mathbf{w}^T \mathbf{x} - \beta = \begin{cases} > 0 & \text{if } \mathbf{x} \text{ belongs to class } o \\ < 0 & \text{if } \mathbf{x} \text{ belongs to class } \Delta \end{cases} \quad (2)$$




Linear Classification



$$\mathbf{w}^T \mathbf{x} - \beta = \begin{cases} > 0 & \text{if } \mathbf{x} \text{ belongs to } \circ \\ < 0 & \text{if } \mathbf{x} \text{ belongs to } \triangle \end{cases}$$



Linear Classification



$$\mathbf{w}^T \mathbf{x} - \beta = \begin{cases} > 0 & \text{if } \mathbf{x} \text{ belongs to } o \\ < 0 & \text{if } \mathbf{x} \text{ belongs to } \Delta \end{cases}$$

The *offset* β can be included in \mathbf{w}


$$\tilde{\mathbf{x}} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix} \quad \tilde{\mathbf{w}} \leftarrow \begin{bmatrix} -\beta \\ \mathbf{w} \end{bmatrix}$$

such that

$$\tilde{\mathbf{w}}^T \tilde{\mathbf{x}} = \mathbf{w}^T \mathbf{x} - \beta.$$



Linear Classification



What is a good w ?

→ We need an **error function** that tells us how good w is.




Two classical Error Functions

Given data $\mathbf{x} \in \mathbb{R}^D$ and corresponding labels $y \in \{-1, +1\}$, two classical error functions $\mathcal{E}(\mathbf{x}, y, \mathbf{w})$ to find the optimal $\mathbf{w} \in \mathbb{R}^D$ are:

Error Function	Used in
$\frac{1}{2}(y - \mathbf{w}^\top \mathbf{x})^2$	Adaline [?]
$\max(0, -y\mathbf{w}^\top \mathbf{x})$	Perceptron [?]



Classification Error as Function of Weights





Given data $\mathbf{x} \in \mathbb{R}^D$ and corresponding labels $y \in \{-1, +1\}$ the classification error \mathcal{E} is a function of the weights \mathbf{w} (and the data \mathbf{x} , y)

$$\mathcal{E}(\mathbf{w}, \mathbf{x}_m, y_m) = - \sum_{m \in \mathcal{M}} \mathbf{w}^T \mathbf{x}_m y_m \quad (3)$$


where \mathcal{M} denotes the index set of all *misclassified* data \mathbf{x}_m



Classification Error as Function of Weights



Gradient Descent




How to minimize the error function?

$$\mathcal{E}(\mathbf{w}, \mathbf{x}_m, y_m) = - \sum_{m \in \mathcal{M}} \mathbf{w}^\top \mathbf{x}_m y_m$$

→ **Gradient Descent**



Gradient Descent



We minimize $\mathcal{E}(\mathbf{w}, \mathbf{x}_m, y_m)$ by walking in the opposite direction of the gradient.

$$\mathbf{w}^{\text{new}} \leftarrow \mathbf{w}^{\text{old}} - \eta \frac{1}{|\mathcal{X}|} \sum_{i=1}^{|\mathcal{X}|} \nabla \mathcal{E}(\mathbf{w}, \mathbf{x}_i, y_i)$$

where \mathcal{X} is the set of data points and η is called a **learning rate**.



Stochastic Gradient Descent

A noisy estimate of


$$\frac{1}{|\mathcal{X}|} \sum_{i=1}^{|\mathcal{X}|} \nabla \mathcal{E}(\mathbf{w}, \mathbf{x}_i, y_i)$$

is obtained by [?]

$$\mathbf{w}^{\text{new}} \leftarrow \mathbf{w}^{\text{old}} - \eta \nabla \mathcal{E}(\mathbf{w}, \mathbf{x}_i, y_i)$$

Note that only \mathbf{w} is stored and only one data point \mathbf{x}_i and label y_i are considered at a time!

→ Scales to large data sets [?]



References

